



Figure 2: Model-generated response rank of the target concept versus the desired level. Point size and density indicate the number of samples at each coordinate. Results shown for Llama-11B with the secondary concept level randomly sampled. For example, “Humor | Persuasiveness” denotes responses generated independently for each humor level (target concept) while persuasiveness is randomly set for each inference.

	Argument Generation			Story Generation			Structured Text Generation		
	Llama-11B	Gemma-12B	Qwen-14B	Llama-11B	Gemma-12B	Qwen-14B	Llama-11B	Gemma-12B	Qwen-14B
$C_a$ (single)	0.76±0.23	0.95±0.07	0.92±0.11	0.81±0.26	0.95±0.06	0.92±0.10	0.73±0.22	0.94±0.12	0.90±0.12
$C_a   C_b$ fixed	0.51±0.41	0.88±0.14	0.88±0.15	0.36±0.45	0.81±0.22	0.90±0.12	0.31±0.50	0.88±0.15	0.84±0.19
$C_a   C_b$ rand	0.54±0.35	0.83±0.20	0.88±0.16	0.33±0.49	0.74±0.25	0.88±0.15	0.17±0.48	0.79±0.21	0.81±0.21
$C_b$ (single)	0.81±0.22	0.98±0.04	0.96±0.05	0.80±0.19	0.97±0.04	0.93±0.10	0.89±0.14	0.99±0.02	0.99±0.03
$C_b   C_a$ fixed	0.58±0.38	0.83±0.19	0.84±0.18	0.59±0.35	0.69±0.34	0.85±0.18	0.56±0.41	0.91±0.15	0.90±0.14
$C_b   C_a$ rand	0.52±0.40	0.76±0.21	0.81±0.21	0.58±0.34	0.70±0.31	0.83±0.20	0.51±0.39	0.79±0.19	0.83±0.19

Table 1: **Humor–persuasiveness.** Spearman correlations for single-concept and dual-concept (fixed / random) across argument, story, and structured text generation.

	Argument Generation			Story Generation			Structured Text Generation		
	Llama-11B	Gemma-12B	Qwen-14B	Llama-11B	Gemma-12B	Qwen-14B	Llama-11B	Gemma-12B	Qwen-14B
$C_a$ (single)	-0.02±0.52	0.52±0.46	0.65±0.30	0.45±0.46	0.92±0.11	0.89±0.12	0.21±0.56	0.15±0.61	0.64±0.21
$C_a   C_b$ fixed	0.02±0.53	0.02±0.56	0.64±0.34	-0.01±0.53	0.35±0.43	0.74±0.26	0.02±0.45	-0.25±0.53	0.39±0.43
$C_a   C_b$ rand	-0.05±0.51	0.12±0.50	0.63±0.32	-0.07±0.50	0.29±0.47	0.64±0.29	0.08±0.40	-0.19±0.45	0.38±0.43
$C_b$ (single)	0.76±0.25	0.95±0.07	0.93±0.10	0.84±0.21	0.98±0.03	0.96±0.07	0.73±0.28	0.97±0.03	0.93±0.09
$C_b   C_a$ fixed	0.76±0.25	0.83±0.19	0.88±0.14	0.71±0.30	0.86±0.17	0.95±0.08	0.45±0.42	0.79±0.31	0.79±0.26
$C_b   C_a$ rand	0.77±0.29	0.80±0.18	0.84±0.15	0.71±0.31	0.72±0.26	0.92±0.08	0.37±0.47	0.63±0.33	0.76±0.28

Table 2: **Clarity–politeness.** Spearman correlations for single-concept and dual-concept (fixed / random) across argument, story, and structured text generation.

	Argument Generation			Story Generation			Structured Text Generation		
	Llama-11B	Gemma-12B	Qwen-14B	Llama-11B	Gemma-12B	Qwen-14B	Llama-11B	Gemma-12B	Qwen-14B
$C_a$ (single)	0.92±0.09	0.98±0.03	0.99±0.02	0.93±0.09	1.00±0.02	0.98±0.05	0.80±0.24	0.93±0.14	0.96±0.07
$C_a   C_b$ fixed	0.56±0.40	0.97±0.05	0.97±0.05	0.77±0.25	0.96±0.07	0.96±0.06	0.42±0.45	0.77±0.33	0.88±0.15
$C_a   C_b$ rand	0.41±0.43	0.92±0.10	0.94±0.08	0.77±0.23	0.96±0.06	0.96±0.05	0.22±0.48	0.71±0.33	0.86±0.17
$C_b$ (single)	0.75±0.32	0.99±0.03	0.98±0.03	0.67±0.33	0.98±0.04	0.97±0.06	0.66±0.32	0.95±0.08	0.87±0.16
$C_b   C_a$ fixed	0.48±0.47	0.90±0.12	0.94±0.08	0.51±0.42	0.93±0.10	0.91±0.10	0.43±0.50	0.72±0.36	0.76±0.29
$C_b   C_a$ rand	0.45±0.44	0.85±0.15	0.93±0.07	0.41±0.46	0.89±0.12	0.89±0.12	0.40±0.51	0.72±0.26	0.75±0.24

Table 3: **Formality–assertiveness.** Spearman correlations for single-concept and dual-concept (fixed / random) across argument, story, and structured text generation.

achieve significantly higher correlations. Politeness follows the standard pattern: high performance for a single concept, but a drop when clarity is introduced. Similarly, in formality–assertiveness, both concepts exhibit consistently high single-concept control (up to 1.00 for Gemma) but degrade under dual-control conditions.

**General trends.** Three broader insights emerge: (i) Qwen-14B and Gemma-12B consistently out-

perform Llama across all settings. This suggests that larger or more instruction-tuned models better preserve disentanglement between stylistic dimensions. (ii) Dual-concept interference remains a central limitation: even when single-concept control is strong, the introduction of a secondary dimension leads to drops in alignment (Figure 2), suggesting weak compositionality of stylistic control. (iii) Task context strongly modulates controllabil-